

Take-Home Assignment — Software Engineering Intern (Web Platform)

Problem: Medicine Directory

Build a Medicine Directory using the [openFDA Drug Label API](https://open.fda.gov/apis/drug/label/) (public, no API key needed).

1. **Homepage** — search for a medicine by name.

2. **Medicine Details Page** — an SEO-friendly page per medicine (e.g. `/medicine/[name]`), showing active ingredients, purpose, warnings, and dosage.

"SEO-friendly" isn't defined further. Decide what that means, and design for the Indian market — who's landing on this page, what device they're on, what layout builds trust, what should actually rank. We want to hear your reasoning.

Test query:

`https://api.fda.gov/drug/label.json?search=openfda.brand_name:"advil"&limit=5`

API docs: <https://open.fda.gov/apis/drug/label/>

What We're Evaluating

1. **Product thinking** — deliberate calls on layout, content, and SEO for the Indian context, not a generic template. How you handle multiple formulations of the same drug. Whether you show a disclaimer for health info.
2. **Edge cases** — zero results, ambiguous multi-match searches, missing fields, slow requests, malformed responses.
3. **Code quality** — no duplicated logic, clean components, Tailwind used properly, not just made to work.
4. **Your thinking process**, visible in your prompt history (below) — what you researched, what you rejected and why, how fast you iterated.

No one nails all four. We're looking for where you spike, and how much care went in.

Requirements

- **Stack:** React.js or Next.js, with TypeScript. Tailwind required. TanStack Query is a bonus, not a requirement.
- Handle these visibly in the UI, not just in a try/catch:
 - No results found
 - API error or timeout
 - Missing fields on a result
- No backend required.

AI Tools: Mandatory Harness + Full Prompt History

AI use is fine — expected, given the 24-hour window. But it must go through one of these:

- Claude Code
- OpenCode
- Roo Code
- Cursor
- VS Code Copilot (Chat/Agent mode)
- Windsurf/Devin

These tools keep a session log we can pull. We're not penalizing AI use — we want to see how you think: what you asked, what you rejected, how you iterated. That's better signal than the final diff alone.

Submit your full raw prompt history — every prompt, in order. Not a summary.

Example format (illustrative, don't reuse):

None

[Turn 1] "Set up a Next.js app with TypeScript and Tailwind, App Router"

[Turn 2] "Add a search bar on the homepage that calls the openFDA label API"

[Turn 3] "This is returning weird nested arrays for some drugs, why"

[Turn 4] "That fallback UI looks off, redo it without a full-page spinner"

Export tips (exact steps vary by version — check your tool's docs if these don't match):

- **Claude Code:** session logs live on disk, typically under a project-local or `~/ .claude` history folder. Include the transcript, or use your version's export/transcript command.
- **Cursor:** open the chat/Composer history for the session, export or copy the full thread.

- **VS Code Copilot Chat:** open the Chat panel history, export or copy it.
- **Windsurf, Roo Code, OpenCode, Devin:** each has a session history in its chat panel — find the export option. No clean export? Keep a running `.txt/.md` log as you go instead.

Didn't use AI? Just say so in your README.

What to Submit

- Public GitHub repo link (or a zip).
- `prompt-history.md` with your full raw AI history, if used.
- Short `README.md`:
 - Your product/SEO/layout decisions, especially anything India-specific.
 - What you'd do differently with more time.
 - Time spent, and which AI harness (if any) you used.

Timeline

24 hours from receipt.

Note

No hidden trick, no single correct UI. A smaller feature set built with real thought — and a prompt history that shows real thinking — beats every API field dumped on screen.